

# Filesystems

---

Sumner Evans and Sam Sartor

November 9, 2017

Mines Linux Users Group

# Introduction

---

# What are Filesystems?

- Filesystems manage the storage and retrieval of files from storage media.
- Filesystems are an abstraction layer between storage media (SSDs, HDDs, disk drives, even tape drives).
- Filesystems exist on *partitions*, physically contiguous segments of the disk.

## Filesystems are Responsible for...

- **Space management:** filesystems allocate and manage space in discrete chunks. Filesystems must keep track of what data is stored at each chunk.
- **Filenames:** identify a storage location in the file system. Can be case sensitive (ext4) or case insensitive (HFS, NTFS).
- **Directories (folders):** group files into separate collections. Modern filesystems allow arbitrary nesting of directories.
- **Metadata:** filesystems store book-keeping information about their contents (e.g. file sizes, last accessed date, owner and permissions, etc.).
- **Access Control:** prevent unauthorized access to files on disk.
- **Data Integrity:** filesystems must be resilient to failure, some are better at this than others.

# Current Filesystems

---

# Linux

---



# Windows & mac

---





# HFS and HFS+

Apple has made a ton of filesystems with varying degrees of terribleness.

- **HFS:** Hierarchical File System — Introduced in 1985 with the first Apple computer with a hard drive. Had a limitation of 65,535 files and every file had to take up at least  $1 / 65,535$ th of the disk.
- **HFS+:** Released in 1998 to fix some of the issues with HFS. the core of the filesystem uses case-insensitive NFD Unicode strings, which led Linus Torvalds to say that “HFS+ is probably the worst file-system ever”.

**APFS:** Apple Filesystem — Introduced in June 2016 to replace HFS+ and is optimized for SSDs. It fixes some of the problems of HFS+. Basically it replicates the work of other modern filesystems which are actually maintained by large communities.

Apple forcibly upgraded all computers to APFS in macOS High Sierra.

# Flashdrives

---



## Other Options

---

# Alternative Filesystems

---

**B-tree** file system (Btrfs) pronounced “Butter FS” or “better FS” or “b-tree FS” was developed starting in 2007 by Oracle.

## Pros

- Copy on Write
- Mostly self-healing
- Can convert from ext\* to Btrfs

## Cons

- It's being deprecated by Oracle. RHEL 7.4 includes it, but they are transitioning away. The SUSE project will still use and maintain it.





TFS is a work-in-progress filesystem for the Redox operating system. Intended as a modern alternative to ZFS, the feature list is mouth-watering.

## Pros

- Concurrent & non-blocking
- Lightweight full-disk compression
- Zero-overhead revision history
- Automatic corruption detection
- $O(1)$  recursive directory copies
- Designed for solid state drives
- Perfectly resilient to power loss

## Cons

- Not fully implemented yet

# Network Filesystems

---

# What is a network filesystem?

You can access remote storage devices over the internet using a *network filesystem*.





# Virtual Filesystems

---

# What is a virtual filesystem?

A *virtual filesystem* is an abstraction layer which converts some other source of data to a filesystem-like structure.

Colloquially, any filesystem that is not stored on disk is called a virtual filesystem. These can be procedural or abstract other kinds of devices.



A \*Nix tmpfs is a filesystem stored in RAM. They appear as mounted filesystems, but are stored in volatile memory. By default, the /tmp directory is a tmpfs.

## Pros

- Useful for storing temporary files such as downloads and temporary files for programs.
- Saves unnecessary disk I/O.
- Security: you can use it to temporarily store sensitive, decrypted files so that the decrypted data is never written to disk.

## Cons

- Everything in a tmpfs will be deleted on reboot since RAM is volatile.

In \*Nix, everything is a file. This includes things such as process information. The way to access this data is through the `proc` filesystem which provides a convenient and standardized method for dynamically accessing process data held in the kernel.

In Linux, `procfs` contains more than just process information including memory information, network utilization statistics, etc.

Filesystem in Userspace (FUSE) is an interface for creating filesystems without writing any kernel-level code, which makes it incredibly useful for creating virtual filesystems. It is available in Linux, FreeBSD, OpenBSD, NetBSD, OpenSolaris, Minix 3, Android, and macOS.

`libfuse...`

- is a C library
- provides a high-level interface for FUSE
- makes creating new filesystems really easy
- has bindings for Python, Rust, etc.

sshfs is a network filesystem implemented through libfuse. You can use it to mount remote directories via ssh.

## Pros

- Very easy & quick to setup (one command)
- Remote machine only needs ssh installed

## Cons

- Intended to be temporary
- Generally slower than NFS

## Configuration/maintenance

---



**Questions?**

# References


- [https://en.wikipedia.org/wiki/File\\_system](https://en.wikipedia.org/wiki/File_system)
- <http://www.tldp.org/LDP/sag/html/filesystems.html>
- <https://arstechnica.com/gadgets/2008/03/past-present-future-file-systems/2/>



# Copyright Notice

This presentation was from the **Mines Linux Users Group**. A mostly-complete archive of our presentations can be found online at <https://lug.mines.edu>.

Individual authors may have certain copyright or licensing restrictions on their presentations. Please be certain to contact the original author to obtain permission to reuse or distribute these slides.



Colorado School of Mines  
Linux Users Group